

CP for DSAI

Lab 02 / Logistic Regression - Binary Classification

Name = "Muhammad Omer Farooq Bhatti"

Id = "st122498"

The function $g(\theta^T \mathbf{x})$ is a squashing or activation function and is required for Binary Classification task.

$$g(\theta^T \mathbf{x}) = [0, 1]$$

$$\mathbf{h} = g(\theta^T \mathbf{x}) = \frac{1}{1 + e^{-\theta^T \mathbf{x}}}$$

$$\begin{aligned} \frac{dg}{dx} &= \frac{0(1 + e^{-x}) - (-1)(e^{-x})}{(1 + e^{-x})^2} \\ &= \frac{e^{-x}}{(1 + e^{-x})^2} = \frac{e^{-x} + 1 - 1}{(1 + e^{-x})^2} \\ &= \frac{1}{(1 + e^{-x})} - \frac{1}{(1 + e^{-x})^2} \\ &= \frac{1}{(1 + e^{-x})} \left(1 - \frac{1}{(1 + e^{-x})}\right) \\ &= g(1 - g) \end{aligned}$$

Binary Classification

Implementation steps:

1. Prepare your data

- add intercept
- \mathbf{X} and \mathbf{y} and \mathbf{w} in the right shape
 - $\mathbf{X} \rightarrow (m, n)$
 - $\mathbf{y} \rightarrow (m,)$
 - $\mathbf{w} \rightarrow (n,)$
 - where m is number of samples
 - where n is number of features
- train-test split
- feature scale
- clean out any missing data
- (optional) feature engineering

2. Predict and calculate the loss

- The loss function is the *cross entropy* defined as

$$J = -\left(\sum_{i=1}^m y^{(i)} \log(h) + (1 - y^{(i)}) \log(1 - h)\right)$$

where h is defined as the sigmoid function as

$$h = \frac{1}{1 + e^{-\theta^T \mathbf{x}}}$$

3. Calculate the gradient based on the loss

- The gradient of θ_j is defined as

$$\frac{\partial J}{\partial \theta_j} = \sum_{i=1}^m (h^{(i)} - y^{(i)}) x_j$$

- This can be derived by knowing that

$$J = -(y \log h + (1 - y) \log(1 - h))$$

$$h = \frac{1}{1 + e^{-g}}$$

$$g = \theta^T x$$

- Thus, gradient of J in respect to some θ_j is

$$\frac{\partial J}{\partial \theta_j} = \frac{\partial J}{\partial h} \frac{\partial h}{\partial g} \frac{\partial g}{\partial \theta_j}$$

where

$$\frac{\partial J}{\partial h} = \frac{h - y}{h(1 - h)}$$

$$\frac{\partial h}{\partial g} = h(1 - h)$$

$$\frac{\partial g}{\partial \theta_j} = x_j$$

- Thus,

$$\begin{aligned}\frac{\partial J}{\partial \theta_j} &= \frac{\partial J}{\partial h} \frac{\partial h}{\partial g} \frac{\partial g}{\partial \theta_j} \\ &= \frac{h - y}{h(1 - h)} * h(1 - h) * x_j \\ &= (h - y)x_j\end{aligned}$$

4. Update the theta with this update rule

$$\theta_j := \theta_j - \alpha * \frac{\partial J}{\partial \theta_j}$$

where α is a typical learning rate range between 0 and 1

5. Loop 2-4 until max_iter is reached, or the difference between old loss and new loss are smaller than some predefined threshold tolerance
 Getting data `sklearn.datasets.make_classification()` - Generates a random n-class classification problem. https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_classification.html

Step #2 - Fit the Algorithm

$$J = -\left(\sum_{i=1}^m y^{(i)} \log(h) + (1 - y^{(i)}) \log(1 - h)\right)$$

$$J = -(y \log h + (1 - y) \log(1 - h))$$

$$h = \frac{1}{1 + e^{-g}}$$

$$g = \theta^T x$$

$$\begin{aligned} \text{Gradient of cost function}, \frac{\partial J}{\partial \theta_j} &= \frac{\partial J}{\partial h} \frac{\partial h}{\partial g} \frac{\partial g}{\partial \theta_j} \\ &= \frac{h - y}{h(1 - h)} * h(1 - h) * x_j \\ &= (h - y)x_j \end{aligned}$$

$$\theta_j := \theta_j - \alpha * \frac{\partial J}{\partial \theta_j}$$

In [7]:

```
%reset
```

Once deleted, variables cannot be recovered. Proceed (y/[n])? y

In [8]:

```
from sklearn import linear_model
from sklearn.datasets import make_classification, make_blobs
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import numpy as np
import matplotlib.pyplot as plt
```

In [9]:

```
class logisticRegression:
    def __init__(self, X, y, l_rate=0.004, max_iterations=1000, scaling=True, plot=True, method = 'minibatch'):
        if (scaling):
            X=featureScale(X)
        #adding intercepts column in X matrix
        self.X_train=self.add_intercept_column(X)
        self.y_train=y
        self.plot=plot
        self.theta=[]
        self.max_iter=max_iterations
        self.alpha = l_rate

    def fit(self):
        print("The data loaded is of shape {0}x{1}.".format(
            self.X_train.shape[0], self.X_train.shape[1]))
        print("The number of features is {0}.".format(
            self.X_train.shape[1]-1))
        print("The number of samples are {0}.".format(
            self.X_train.shape[0]))
        print("Starting Mini-Batch Gradient Descent on Feature Vector X of size {0}x{1}...".format(
            self.X_train.shape[0], self.X_train.shape[1]))
        self.theta, n_iter = self.mini_batch_GD(
            self.X_train, self.y_train, self.max_iter, self.plot)
        print("Finished Mini-Batch Gradient Descent on Feature Vector X of size {0}x{1} in {2} iterations.".format(
            self.X_train.shape[0], self.X_train.shape[1], n_iter))

    def featureScale(self,_X):
        scaler=StandardScaler()
        _X=scaler.fit_transform(_X)
        return _X

    def add_intercept_column(self,_X):
        #Concatenating the intercept to training set as variable b of shape = (no. of samples used for training, 1)
        b=np.ones((_X.shape[0],1))
        _X = np.concatenate((b, _X), axis=1)

        return _X

    def mini_batch_GD(self, X,y, max_iter, plot):
        #Using mini-batch Gradient Descent with replacement
        _theta=np.zeros(X.shape[1])
        batch_size = int(0.1*(X.shape[0]))
        #alpha = 0.001
        J_iter=np.zeros(max_iter)
        for i in range(max_iter):

            index=np.random.randint(0,X.shape[0]) #choose a contiguous block of samples of size = batch_size at random
            xtrain=X[index:index+batch_size]
            ytrain=y[index:index+batch_size]

            cost, grad = self.gradient(_theta, X, y)
            J_iter[i]= cost
            _theta = _theta - self.alpha * grad

        if (plot==True):
            self.plot_J_iter(J_iter,i)
        return _theta,i

    def plot_J_iter(self, J_iter, n_iter):
        plt.plot(np.arange(1,n_iter,1), J_iter[0:n_iter-1])
        plt.xlabel("Number of iterations")
        plt.ylabel("Loss")
        plt.title("Loss function vs No. of Iterations ")

    def sigmoid(self, X, deriv = False):
        g = 1/(1+np.exp(-X))      #Where X = X@theta
        if deriv==True:
            g_deriv = g*(1-g)     #It may be useful to plot the derivative function as value of alpha (learning rate)
            return g_deriv         #is limited by the maximum value of the derivative function
        else:
            return g

    def h_theta(self, X, _theta):
        return self.sigmoid(np.dot(X,_theta))

    def gradient(self, _theta, X, y):
        h = self.h_theta(X,_theta)
        error = h-y
        grad = np.dot(X.T, error) #(X(mxn) dot error(mx1) not right as (n!=m) so we do X.T(nxm) dot error(mx1) = grad(nx1)
        cost = -np.sum( y*np.log(h) + (1-y)*np.log(1-h) )
        return cost, grad

    def predict(self, xtest):
        ##Concatenating the intercept to testing set as variable b of shape = (no. of samples used for testing, 1)
        #adding intercepts column in X matrix
        xtest = self.add_intercept_column(xtest)

        return np.round(self.h_theta(xtest, self.theta))
```

```
In [10]: class classificationReport:
    def __init__(self, y_test, y_hat):
        self.ytest = y_test
        self.yhat = y_hat

    for (p,n) in [(0,1), (1,0)]:
        print("-----\n")
        print(f"Computing for Positive = {p}, Negative = {n}\n")
        self.c_matrix, n_TP, n_FP, n_FN, n_TN = self.confusion_matrix(self.ytest, self.yhat, p, n)
        self.accuracy = self.calc_accuracy(n_TP, n_TN, n_FP, n_FN)
        self.precision = self.calc_precision(n_TP, n_FP)
        self.recall = self.calc_recall(n_TP, n_FN)
        self.f1 = self.calc_f1()
        print(f"Accuracy: {self.accuracy}")
        print(f"Precision: {self.precision}")
        print(f"Recall: {self.recall}")
        print(f"F1: {self.f1}")
        print("\n-----\n")

    def calc_accuracy(self, n_TP, n_TN, n_FP, n_FN):
        return (n_TP + n_TN) / (n_TP + n_TN + n_FP + n_FN)

    def calc_precision(self, n_TP, n_FP):
        return (n_TP) / (n_TP + n_FP)

    def calc_recall(self, n_TP, n_FN):
        return (n_TP) / (n_TP + n_FN)


    def calc_f1(self):
        return 2 * (self.precision * self.recall) / (self.precision + self.recall)

    def confusion_matrix(self, ytest, yhat, p, n):
        c_matrix = []
        n_TP=0
        n_FP=0
        n_FN=0
        n_TN=0

        for i in range(len(ytest)):
            if ( (ytest[i]==p) and (yhat[i] == p) ):
                c_matrix.append("TP")
                n_TP+=1
            elif ( (yhat[i]==p) and (ytest[i]==n) ):
                c_matrix.append("FP")
                n_FP+=1
            elif ( (yhat[i]==n) and (ytest[i]==p) ):
                c_matrix.append("FN")
                n_FN+=1
            elif ( (yhat[i]==n) and (ytest[i]==n) ):
                c_matrix.append("TN")
                n_TN+=1
        return c_matrix, n_TP, n_FP, n_FN, n_TN
```

```
In [26]: X, y = make_blobs(n_samples=500, centers=2, n_features=2,
                      random_state=0)

plt.scatter(X[:, 0], X[:, 1], marker='o', c=y,
            s=25, edgecolor='k')
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3)
```




```
In [27]: model = logisticRegression(X_train,y_train, l_rate = 0.0001, max_iterations = 5000)
model.fit()
yhat = model.predict(X_test)
print(f"The computed parameter space, theta = {model.theta}")
report = classificationReport(y_test, yhat)
```

The data loaded is of shape 350x3.
The number of features is 2.
The number of samples are 350
Starting Mini-Batch Gradient Descent on Feature Vector X of size 350x3...
Finished Mini-Batch Gradient Descent on Feature Vector X of size 350x3 in 4999 iterations.

The computed parameter space, theta = [0.24669521 0.95217017 -5.09491161]

Computing for Positive = 0, Negative = 1
Accuracy: 0.7133333333333334
Precision: 0.656
Recall: 1.0
F1: 0.7922705314009661

```
Computing for Positive = 1, Negative = 0
Accuracy: 0.7133333333333334
Precision: 1.0
Recall: 0.36764705882352944
F1: 0.5376344086021506
```



```
In [28]: from sklearn.metrics import classification_report
print("=====Classification report=====")
print(classification_report(y_test, yhat))
```

```
=====Classification report=====
      precision    recall  f1-score   support

          0       0.66     1.00     0.79      82
          1       1.00     0.37     0.54      68

   accuracy                           0.71      150
  macro avg       0.83     0.68     0.66      150
weighted avg       0.81     0.71     0.68      150
```

```
In [ ]:
```